


```
In [1]: #1 IONOSPHERE
import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv(r"C:\Users\venka\OneDrive\Documents\ionosphere_data karthik.csv")
df
```

Out[2]:

	column_a	column_b	column_c	column_d	column_e	column_f	column_g	column_h	co
0	True	False	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1
1	True	False	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1
2	True	False	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	(
3	True	False	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	(
4	True	False	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	(
...
346	True	False	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	(
347	True	False	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	(
348	True	False	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	(
349	True	False	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	(
350	True	False	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	(

351 rows × 35 columns



```
In [3]: pd.set_option('display.max_rows',10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
print('This dataframe has %d rows and %d columns'%(df.shape))
```

This dataframe has 351 rows and 35 columns

In [4]: `df.head(10)`

Out[4]:

	column_a	column_b	column_c	column_d	column_e	column_f	column_g	column_h	column_i
0	True	False	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.0
1	True	False	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.0
2	True	False	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.8
3	True	False	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.0
4	True	False	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.7
5	True	False	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.1
6	True	False	0.97588	-0.10602	0.94601	-0.20800	0.92806	-0.28350	0.8
7	False	False	0.00000	0.00000	0.00000	0.00000	1.00000	-1.00000	0.0
8	True	False	0.96355	-0.07198	1.00000	-0.14333	1.00000	-0.21313	1.0
9	True	False	-0.01864	-0.08459	0.00000	0.00000	0.00000	0.00000	0.1

In [5]: `features_matrix=df.iloc[:,0:34]`

In [6]: `target_vector=df.iloc[:,-1]`

In [7]: `print('The features matrix has %d rows and %d columns'%(features_matrix.shape`

The features matrix has 351 rows and 34 columns

In [8]: `print('The target matrix has %d rows and %d columns'%(np.array(target_vector)`


The target matrix has 351 rows and 1 columns

In [9]: `features_matrix_standardized=StandardScaler().fit_transform(features_matrix)`

In [10]: `algorithm=LogisticRegression(penalty='l2',dual=False,tol=1e-4,C=1.0,fit_inter`

In [21]: `Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_v`

In [22]: `observation=[[1,0,0.99539,-0.05889,0.8524299999999999,0.02306,0.8339799999999999,`



In [23]: `from sklearn.linear_model import LogisticRegression
predictions=Logistic_Regression_Model.predict(observation)`

In [24]: `print('The model predicted the observation to belong to class %s'%(prediction`

The model predicted the observation to belong to class ['g']

In [25]: `print('The algorithm was trained to predict one of the two classes :%s'%(algo`

The algorithm was trained to predict one of the two classes :['b' 'g']

In [26]: `print("""The model says the probability of the observation we passed belonging`



The model says the probability of the observation we passed belonging to class ['b'] is 0.00777393160013784

In [27]: `print()`

In [28]: `print("""The model says the probability of the observation we passed belonging`



The model says the probability of the observation we passed belonging to class ['g'] is 0.9922260683998622

In [29]: `#2 DIGITS
import re
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
%matplotlib inline
digits=load_digits()`


```
In [21]: print("Image Data Shape",digits.data.shape)
print("Label Data Shape",digits.target.shape)
```

Image Data Shape (1797, 64)
Label Data Shape (1797,)

```
In [22]: plt.figure(figsize=(20,4))
```

Out[22]: <Figure size 2000x400 with 0 Axes>
<Figure size 2000x400 with 0 Axes>

```
In [23]: for index,(image,label)in enumerate(zip(digits.data[0:5],digits.target[0:5])):
plt.subplot(1,5,index+1)
plt.imshow(np.reshape(image,(8,8)),cmap=plt.cm.gray)
plt.title('Training:%i\n'%label,fontsize=10)
```



```
In [24]: x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test
```

```
In [25]: print(x_train.shape)
```

(1257, 64)

```
In [26]: print(y_train.shape)
```

(1257,)

```
In [27]: print(x_test.shape)
```

(540, 64)

In [28]: `print(y_test.shape)`

(540,)

In [29]: `from sklearn.linear_model import LogisticRegression`
`logisticRegr=LogisticRegression(max_iter=10000)`

In [30]: `logisticRegr.fit(x_train,y_train)`

Out[30]:

LogisticRegression

LogisticRegression(max_iter=10000)

In [31]: `print(logisticRegr.predict(x_test))`

```
[4 0 9 1 8 7 1 5 1 6 6 7 6 1 5 5 8 6 2 7 4 6 4 1 5 2 9 5 4 6 5 6 3 4 0 9 9
 8 4 6 8 8 5 7 9 8 9 6 1 7 0 1 9 7 3 3 1 8 8 8 9 8 5 8 4 9 3 5 8 4 3 1 3 8
 7 3 3 0 8 7 2 8 5 3 8 7 6 4 6 2 2 0 1 1 5 3 5 7 1 8 2 2 6 4 6 7 3 7 3 9 4
 7 0 3 5 1 5 0 3 9 2 7 3 2 0 8 1 9 2 1 5 1 0 3 4 3 0 8 3 2 2 7 3 1 6 7 2 8
 3 1 1 6 4 8 2 1 8 4 1 3 1 1 9 5 4 8 7 4 8 9 5 7 6 9 4 0 4 0 0 9 0 6 5 8 8
 3 7 9 2 0 8 2 7 3 0 2 1 9 2 7 0 6 9 3 1 1 3 5 2 5 5 2 1 2 9 4 6 5 5 5 9 7
 1 5 9 6 3 7 1 7 5 1 7 2 7 5 5 4 8 6 6 2 8 7 3 7 8 0 9 5 7 4 3 4 1 0 3 3 5
 4 1 3 1 2 5 1 4 0 3 1 5 5 7 4 0 1 0 9 5 5 5 4 0 1 8 6 2 1 1 1 7 9 6 7 9 7
 0 4 9 6 9 2 7 2 1 0 8 2 8 6 5 7 8 4 5 7 8 6 4 2 6 9 3 0 0 8 0 6 6 7 1 4 5
 6 9 7 2 8 5 1 2 4 1 8 8 7 6 0 8 0 6 1 5 7 8 0 4 1 4 5 9 2 2 3 9 1 3 9 3 2
 8 0 6 5 6 2 5 2 3 2 6 1 0 7 6 0 6 2 7 0 3 2 4 2 3 6 9 7 7 0 3 5 4 1 2 2 1
 2 7 7 0 4 9 8 5 6 1 6 5 2 0 8 2 4 3 3 2 9 3 8 9 9 5 9 0 3 4 7 9 8 5 7 5 0
 5 3 5 0 2 7 3 0 4 3 6 6 1 9 6 3 4 6 4 6 7 2 7 6 3 0 3 0 1 3 6 1 0 4 3 8 4
 3 3 4 8 6 9 6 3 3 0 5 7 8 9 1 5 3 2 5 1 7 6 0 6 9 5 2 4 4 7 2 0 5 6 2 0 8
 4 4 4 7 1 0 4 1 9 2 1 3 0 5 3 9 8 2 6 0 0 4]
```

In [32]: `score=logisticRegr.score(x_test,y_test)`
`print(score)`

0.9537037037037037

In []:

```
In [1]: #3 GENDER SUBMISSION
import re
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import metrics
%matplotlib inline
```

```
In [6]: df=pd.read_csv(r"C:\Users\venka\OneDrive\Documents\gender_submission.csv")
print(df)
```

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1
..
413	1305	0
414	1306	1
415	1307	0
416	1308	0
417	1309	0

[418 rows x 2 columns]

```
In [7]: plt.figure(figsize=(20,2))
```

```
Out[7]: <Figure size 2000x200 with 0 Axes>
```

```
<Figure size 2000x200 with 0 Axes>
```

```
In [8]: df.describe()
```

```
Out[8]:
```

	PassengerId	Survived
count	418.000000	418.000000
mean	1100.500000	0.363636
std	120.810458	0.481622
min	892.000000	0.000000
25%	996.250000	0.000000
50%	1100.500000	0.000000
75%	1204.750000	1.000000
max	1309.000000	1.000000

```
In [9]: df.isnull().any()
```


```
Out[9]: PassengerId    False  
Survived             False  
dtype: bool
```

```
In [10]: pd.set_option('display.max_rows',10000000000)  
pd.set_option('display.max_columns',10000000000)  
pd.set_option('display.width',95)  
print('This dataframe has %d rows and %d columns'%(df.shape))
```

This dataframe has 418 rows and 2 columns

```
In [11]: sns.boxplot(df['PassengerId'])
```

```
Out[11]: <Axes: >
```



```
In [12]: df.head(10)
```

```
Out[12]:
```

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1
5	897	0
6	898	1
7	899	0
8	900	1
9	901	0

```
In [13]: features_matrix=df.iloc[:,0:1]
```

```
In [14]: target_vector=df.iloc[:, -1]
```

```
In [15]: print('The features matrix has %d rows and %d columns'%(features_matrix.shape
```


The features matrix has 418 rows and 1 columns

```
In [16]: print('The target matrix has %d rows and %d columns'%(np.array(target_vector)
```

The target matrix has 418 rows and 1 columns

```
In [17]: from sklearn.preprocessing import StandardScaler  
features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

```
In [18]: algorithm=LogisticRegression(penalty='l2',dual=False,tol=1e-4,C=1.0,fit_inter
```



```
In [21]: Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_v
```



```
In [22]: observation=[[1]]
```

```
In [27]: predictions=Logistic_Regression_Model.predict(observation)
```

```
In [33]: print('The model predicted the observation to belong to class %s'%(prediction))
```

The model predicted the observation to belong to class ['g']

```
In [34]: print('The algorithm was trained to predict one of the two classes :%s' %(algorithm))
```

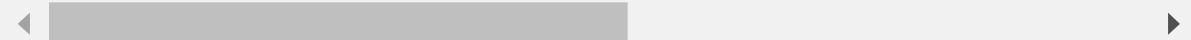
The algorithm was trained to predict one of the two classes :['b' 'g']

```
In [30]: print("""The model says the probability of the observation we passed belonging to class ['0'] is 0.6474324251144166""")
```



The model says the probability of the observation we passed belonging to class ['0'] is 0.6474324251144166

```
In [31]: print("""The model says the probability of the observation we passed belonging to class ['1'] is 0.35256757488558343""")
```



The model says the probability of the observation we passed belonging to class ['1'] is 0.35256757488558343

```
In [ ]:
```